

# Lecture 18: Classification

Big Data and Machine Learning for Applied Economics  
Econ 4676

Ignacio Sarmiento-Barbieri

Universidad de los Andes

October 19, 2021

# Agenda

- 1 Recap: Regularization and Lasso for Causality
- 2 Classification
  - K-Nearest Neighbors
  - Logit
  - Linear Discriminant Analysis
- 3 Misclassification Rates
  - ROC curve
- 4 Review & Next Steps
- 5 Further Readings
- 6 Demos in R
  - KNN
  - Logit
  - LDA
  - ROC

# Elastic Net

- ▶ Naive Elastic Net

$$\min_{\beta} NEL(\beta) = \sum_{i=1}^n (y_i - x_i' \beta)^2 + \lambda_1 \sum_{s=2}^p |\beta_s| + \lambda_2 \sum_{s=2}^p \beta_s^2 \quad (1)$$

- ▶ Elastic Net: rescaled version
- ▶ Double Shrinkage introduces “too” much bias, *final* version “corrects” for this

$$\hat{\beta}_{EN} = \frac{1}{\sqrt{1 + \lambda_2}} \hat{\beta}_{naive EN} \quad (2)$$

- ▶ Careful sometimes software asks.

# Lasso for Causality

## Inference with Selection among Many Controls

$$y_i = \alpha D_i + X_i' \theta_y + r_{yi} + \zeta_i \quad (3)$$

- ▶ We apply variable selection methods to each of the two reduced form equations and then use all of the selected controls in estimation of  $\alpha$ .
- ▶ We select
  - 1 A set of variables that are useful for predicting  $y_i$ , say  $X_{yi}$ , and
  - 2 A set of variables that are useful for predicting  $D_i$ , say  $X_{di}$ .
- ▶ We then estimate  $\alpha$  by ordinary least squares regression of  $y_i$  on  $d_i$  and the union of the variables selected for predicting  $y_i$  and  $D_i$ , contained in  $X_{yi}$  and  $X_{di}$ .
- ▶ We thus make sure we use variables that are important for either of the two predictive relationships to guard against OVB

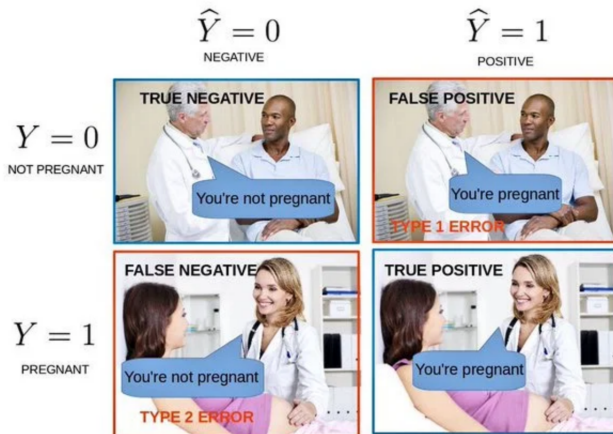
## Classification

# Classification: Motivation

- ▶ Admit a student to *PEG* based on their grades and LoR
- ▶ Give a credit, based on credit history, demographics?
- ▶ Classifying emails: spam, personal, social based on email contents
- ▶ Aim is to classify  $y$  based on  $X$ 's
- ▶  $y$  can be
  - ▶ qualitative (e.g., spam, personal, social)
  - ▶ Not necessarily ordered
  - ▶ Not necessarily two categories, but will start with the binary case

# Motivation

- ▶ Two states of nature  $y \rightarrow n \in \{0,1\}$
- ▶ Two actions  $(\hat{y}) \rightarrow a \in \{0,1\}$



Source: <https://dzone.com/articles/understanding-the-confusion-matrix>

# Probability, Cost, and Classification

- ▶ Two states of nature  $y \rightarrow n \in \{0, 1\}$
- ▶ Two actions  $(\hat{y}) \rightarrow a \in \{0, 1\}$
- ▶ Probabilities
  - ▶  $p = Pr(y = 1|X)$
  - ▶  $1 - p = Pr(y = 0|X)$
- ▶ Loss:  $L(a, )$ , penalizes being in bin  $(a, n)$
- ▶ Risk: expected loss of taking action  $a$



# Probability, Cost, and Classification

- Risk: expected loss of taking action  $a$

$$E[L(a, n)] = \sum_n p_n L(a, n) \quad (4)$$

$$R(a) = (1 - p)L(a, 0) + pL(a, 1)$$

- The objective is the same as before: minimize the risk
- We have to define  $L(a, n)$

# Probability, Cost, and Classification

- ▶ Which action do we choose?

# Probability, Cost, and Classification

- ▶ Which action do we choose?
- ▶ We can compare the risk of each action
- ▶ We are going to choose to take action 1 when the risk is lower:

$$\begin{aligned} R(1) &< R(0) \\ 1 - p &< p \\ p &> \frac{1}{2} \end{aligned} \tag{5}$$

- ▶ This is known as the Bayes Classifier, choose the estate that minimizes the risk


# Probability, Cost, and Classification

- ▶ Under a 0-1 penalty the problem boils down to finding  $p = Pr(y = 1|X)$
- ▶ We then predict 1 if  $p > 0.5$  and 0 otherwise (Bayes classifier)
- ▶ We can think 3 ways of finding this probability in binary cases
  - ▶ K-Nearest Neighbors
  - ▶ Logistic
  - ▶ LDA
- ▶ Why not  $p = X\beta$ ?

## K-Nearest Neighbors

# K-Nearest Neighbors

- K nearest neighbor (K-NN) algorithm predicts class  $\hat{y}$  for  $x$  by asking *What is the most common class for observations around  $x$ ?*



Source: Taddy (2019)

Lecture 18

# K-Nearest Neighbors

- ▶ K nearest neighbor (K-NN) algorithm predicts class  $\hat{y}$  for  $x$  by asking *What is the most common class for observations around  $x$ ?*
- ▶ Algorithm: given an input vector  $x_f$  where you would like to predict the class label
  - ▶ Find the K nearest neighbors in the dataset of labeled observations,  $x_i, y_{i=1}^n$ , the most common distance is the Euclidean distance (units):

$$d(x_i, x_f) = \sqrt{\sum_{j=1}^p (x_{ij} - x_{fj})^2} \quad (6)$$

- ▶ This yields a set of the  $K$  nearest observations with labels:

$$[x_{i1}, y_{i1}], \dots, [x_{iK}, y_{iK}] \quad (7)$$

- ▶ The predicted class of  $x_f$  is the most common class in this set

$$\hat{y}_f = \text{mode}\{y_{i1}, \dots, y_{iK}\} \quad (8)$$

# K-Nearest Neighbors


- ▶ There are some major problems with practical applications
  - ▶ Knn predictions are unstable as a function of  $K$

$$K = 1 \implies \hat{p}(\text{white}) = 0$$

$$K = 2 \implies \hat{p}(\text{white}) = 1/2$$

$$K = 3 \implies \hat{p}(\text{white}) = 2/3$$

$$K = 4 \implies \hat{p}(\text{white}) = 1/2$$



Source: Taddy (2019)



# K-Nearest Neighbors

- ▶ In this case
  - ▶ 1-Knn manages 70% accuracy
  - ▶ 5-Knn manages 60% accuracy
  - ▶ H.W. try for different seed, (Taddy's is 80% and 70%)
- ▶ There are some major problems with practical implications
  - ▶ Knn predictions are unstable as a function of  $K$
  - ▶ This instability of prediction makes it hard to choose the optimal  $K$  and cross validation doesn't work well for KNN
  - ▶ Since prediction for each new  $x$  requires a computationally intensive counting, KNN is too expensive to be useful in most big data settings.
  - ▶ KNN is a good idea, but too crude to be useful in practice

## Logit

# Logit

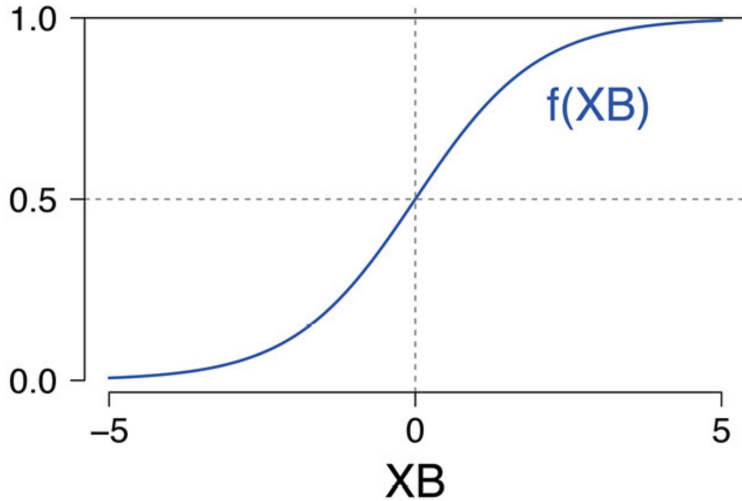
We have a conditional probability

$$\Pr(y = 1|X) = f(X'\beta) \quad (9)$$

Logistic regression uses a *logit* (sigmoid, softmax) link function

$$\log \left( \frac{p(y = 1|X)}{1 - p(y = 1|X)} \right) = \beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k \quad (10)$$

# Logit



Source: Taddy (2019)

# Logit

We have a conditional probability

$$Pr(y = 1|X) = f(X'\beta) \quad (11)$$

Can recover predictions:

$$p(y = 1|X) = \frac{e^{X'\beta}}{1 + e^{X'\beta}} = \frac{\exp(\beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k)}{1 + \exp(\beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k)} \quad (12)$$

## Linear Discriminant Analysis

# Linear Discriminant Analysis

Reverend Bayes to the rescue: Bayes Theorem

$$p(y = 1|X) = \frac{f(X|y = 1)p(y = 1)}{m(X)} \quad (13)$$

with  $m(X)$  is the marginal distribution of  $X$ , i.e.

$$m(X) = \int f(X|y = 1)p(y = 1)dy \quad (14)$$

Recall that there are two states of nature  $y \rightarrow i \in \{0, 1\}$

$$\begin{aligned} m(X) &= f(X|y = 1)p(y = 1) + f(X|y = 0)p(y = 0) \\ &= f(X|y = 1)p(y = 1) + f(X|y = 0)(1 - p(y = 1)) \end{aligned} \quad (15)$$

# Linear Discriminant Analysis

- ▶ This is basically an empirical Bayes approach
- ▶ We need to estimate  $f(X|y = 1), f(X|y = 0)$  and  $p(y = 1)$ 
  - ▶ Let's start by estimating  $p(y = 1)$ . We've done this before

$$p(y = 1) = \frac{\sum_{i=1}^n 1[y_i = 1]}{N} \quad (16)$$

- ▶ Next  $f(X|y = j)$  with  $j = 0, 1$ .
  - ▶ if we assume one predictor and  $X|y \sim N(\mu_j, \sigma_j)$
  - ▶ the problem boils down to estimating  $\mu_j, \sigma_j$
  - ▶ LDA makes it simpler, assumes  $\sigma_j = \sigma \forall j$
  - ▶ then partition the sample in two  $y = 0$  and  $y = 1$ , estimate the moments and get  $\hat{f}(X|y = j)$
- ▶ Plug everything into the Bayes Rule and you're done



# Linear Discriminant Analysis

## Extensions

- ▶ If we have  $k$  predictors?
- ▶ then  $X|y \sim NM(\mu, \Sigma)$

$$f(X|Y = j) = \frac{1}{(2\pi)^{k/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_j)' \Sigma_j (x - \mu_j)\right) \quad (17)$$

- ▶  $\mu_j$  is the vector of the sample means in each partition  $j = 0, 1$
- ▶  $\Sigma_j$  is the matrix of variance and covariances of each partition  $j = 0, 1$
- ▶ Can we lift normality?

# Linear Discriminant Analysis

- ▶ Why is it call linear?
- ▶ Note

$$p > \frac{1}{2} \iff \ln\left(\frac{p}{(1-p)}\right) \quad (18)$$

- ▶ Logit with one predictor

$$\beta_1 + \beta_2 X \quad (19)$$

- ▶ Classification: in the probability of space
- ▶ Discrimination: in the space of  $X$
- ▶  $\beta_1 + \beta_2 X$  is the discrimination function for logit (it is lineal)

# Linear Discriminant Analysis

- ▶ LDA?
- ▶ One predictor with  $\sigma_0 = \sigma_1$  (equal variance)

$$p(Y = 1|X) = \frac{f(X|Y = 1)p(Y = 1)}{f(X|Y = 1)p(Y = 1) + f(X|Y = 0)(1 - p(Y = 1))} \quad (20)$$

- ▶ Then under the equal variance assumption

$$\frac{p(Y = 1|X)}{1 - p(Y = 1|X)} = \frac{f(X|Y = 1)p(Y = 1)}{f(X|Y = 0)(1 - p(Y = 1))} \quad (21)$$

$$= \frac{p(Y = 1)\exp((x - \mu_1)^2)}{(1 - p(Y = 1))\exp((x - \mu_0)^2)} \quad (22)$$

# Linear Discriminant Analysis

- ▶ Taking logs

$$\log \left( \frac{p(Y = 1|X)}{1 - p(Y = 1|X)} \right) = \log \left( \frac{p(Y = 1)}{(1 - p(Y = 1))} + (x - \mu_1)^2 - (x - \mu_0)^2 \right) \quad (23)$$

$$= \log \left( \frac{p(Y = 1)}{(1 - p(Y = 1))} + \mu_1^2 - \mu_0^2 - 2(\mu_1 - \mu_0)x \right) \quad (24)$$


$$= \gamma_1 + \gamma_2 X \quad (25)$$

- ▶ under the assumption of equal variance the discrimination function is lineal
- ▶ Note: logit estimates  $\gamma_1$  and  $\gamma_2$

## Misclassification Rates


# Misclassification Rates

- Predicted probabilities from Logit model



# Misclassification Rates

- ▶ A classification rule, or cutoff, is the probability  $p$  at which you predict
  - ▶  $\hat{y}_i = 0$  if  $p_i < p$
  - ▶  $\hat{y}_i = 1$  if  $p_i \geq p$
- ▶ Measures of performance
  - ▶ *1-Specificity*: False Positive Rate, Type I error
  - ▶ *Sensitivity*: True Positive Rate, power, (1-Type II error)




# ROC

- ▶ ROC curve: Receiver operating characteristic curve
- ▶ ROC curve illustrates the trade-off of the classification rule
- ▶ Gives us the ability
  - ▶ Measure the predictive capacity of our model
  - ▶ Compare between models
- ▶ Some definitions
  - ▶  $P = \sum Y_i$  positives
  - ▶  $N = \sum (1 - Y_i)$  negatives
  - ▶  $T = P + N$  all observations
  - ▶ True Positives:  $TP = \sum \hat{Y}_i Y_i$ , True Positive Rate =  $\frac{TP}{P}$
  - ▶ False Positives:  $FP = \sum \hat{Y}_i (1 - Y_i)$ , False Positive Rate =  $\frac{FP}{N}$



# ROC

- ▶ Binary Classifier:  $\hat{Y}_i = 1[p_i > c], c \in [0, 1]$
- ▶ Bayes fixes  $c = 0.5$
- ▶ Ideally  $TPR = 1$  and  $FPR = 0$
- ▶ ROC curve give us the locus of all possible  $TPR$  and  $FPR$  for all possible  $c \in [0, 1]$



## ► ROC Properties

### ► Has positive slope

► In  $(0,0)$ ,  $c = 1$ . When  $c \downarrow$ ,  $TP \uparrow$  and  $FP \uparrow$ . Then

$$TPR = \sum \frac{\hat{Y}_i Y_i}{P} \quad FPR = \sum \frac{\hat{Y}_i (1 - Y_i)}{T - P} \quad (26)$$

► Is easy to show

$$TPR = \frac{\sum \hat{Y}_i}{P} - \frac{T - P}{P} FPR \quad (27)$$

► ROC is the locus of all possible  $TPR$  and  $FPR$  for all possible  $c \in [0, 1]$

$$TPR = \frac{\sum \hat{Y}_i(c)}{P} - \frac{T - P}{P} FPR(c) \quad (28)$$

# ROC: Summary

- ▶ Ideal ROC curve
- ▶ AUC: area under the curve, is like an  $R^2$
- ▶ Help us compare between classifiers
- ▶ Dominated classifiers?
- ▶ Which  $c$ ? Choose a max  $FPR$

# Review & Next Steps

- ▶ Review Classification:
  - ▶ KNN
    - ▶ Intuitive
    - ▶ Not very useful in practice, curse of dimensionality
  - ▶ Logit
  - ▶ Linear Discriminant Analysis
  - ▶ Misclassification Rates: ROC curve
  - ▶ QDA?
  - ▶ Multiple Classes?
- ▶ Next class: Problem Sets, Text Data!
- ▶ Questions? Questions about software?

## Further Readings

- ▶ Friedman, J., Hastie, T., & Tibshirani, R. (2001). The elements of statistical learning (Vol. 1, No. 10). New York: Springer series in statistics.
- ▶ James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). An introduction to statistical learning (Vol. 112, p. 18). New York: springer.
- ▶ Kuhn, M. (2012). The caret package. R Foundation for Statistical Computing, Vienna, Austria. <https://topepo.github.io/caret/index.html>
- ▶ Taddy, M. (2019). Business data science: Combining machine learning and economics to optimize, automate, and accelerate business decisions. McGraw Hill Professional.
- ▶ Zou, H. y Hastie, T., 2005, Regularization and variable selection via the elastic net, Journal of the Royal Statistical Society, 67, 2, 301-320.


# K-Nearest Neighbors

```
#Load the required packages
library("class") #for KNN
library("MASS") # a library of example datasets
#Read the data
data(fgl) ## loads the data into R; see help(fgl)
str(fgl)
```

```
## 'data.frame':    214 obs. of  10 variables:
## $ RI : num  3.01 -0.39 -1.82 -0.34 -0.58 ...
## $ Na : num  13.6 13.9 13.5 13.2 13.3 ...
## $ Mg : num  4.49 3.6 3.55 3.69 3.62 3.61 3.6 3.61 3.58 3.6 ...
## $ Al : num  1.1 1.36 1.54 1.29 1.24 1.62 1.14 1.05 1.37 1.36 ...
## $ Si : num  71.8 72.7 73 72.6 73.1 ...
## $ K : num  0.06 0.48 0.39 0.57 0.55 0.64 0.58 0.57 0.56 0.57 ...
## $ Ca : num  8.75 7.83 7.78 8.22 8.07 8.07 8.17 8.24 8.3 8.4 ...
## $ Ba : num  0 0 0 0 0 0 0 0 0 0 ...
## $ Fe : num  0 0 0 0 0 0.26 0 0 0 0.11 ...
## $ type: Factor w/ 6 levels "WinF","WinNF",...: 1 1 1 1 1 1 1 1 1 1 ...
```

Refractive index and chemical composition for six possible glass types: float glass window (WinF), nonfloat glass window (WinNF), vehicle window (Veh), container (Con), tableware (Tabl), vehicle headlamp (Head)

# K-Nearest Neighbors



# K-Nearest Neighbors

- ▶ Units matter
  - ▶ Since distance is measured on the raw  $x$  values, units matter.
  - ▶ As we did for regularization, we will standardized observations.
  - ▶ R `scale` function does this, i.e., convert columns to mean-zero sd-one

```
x <- scale(fgl[,1:9]) # column 10 is the class label
apply(x,2,sd) # see ?apply
```

```
## RI Na Mg Al Si K Ca Ba Fe
## 1 1 1 1 1 1 1 1 1
```



# K-Nearest Neighbors

## ► Before running Knn

- Make sure you have numeric matrices of training data  $x$  values, with labels  $y$
- Also need to provide new *test* values where you would like to predict
- Note that there's no model to fit, Knn, just counts neighbors for each observation in test

```
set.seed(1010101)
test <- sample(1:214,10)
nearest1 <- knn(train=x[-test,], test=x[test,], cl=fgl$type[-test], k=1)
nearest5 <- knn(train=x[-test,], test=x[test,], cl=fgl$type[-test], k=5)
data.frame(fgl$type[test],nearest1,nearest5)
```

```
##      fgl.type.test. nearest1 nearest5
## 1           WinF         WinF      WinNF
## 2             Head         Head        Head
## 3          WinNF      WinNF      WinNF
## 4           WinF         WinF        WinF
## 5          WinNF      WinNF      WinNF
## 6          WinNF      WinNF      WinNF
## 7           Head         Con         Con
## 8           Head      WinNF      WinNF
## 9          WinNF      WinNF      WinNF
## 10         WinNF         WinF        WinF
```

# Logit Demo

```
set.seed(101010) #sets a seed
credit<-readRDS("credit_class.rds")
#70% train
indic<-sample(1:nrow(credit),floor(.7*nrow(credit)))
#Partition the sample
train<-credit[indic,]
test<-credit[-indic,]
head(credit)
```

##	Default	duration	amount	installment	age	history	purpose	foreign	rent
## 1	0	6	1169		4	67 terrible	goods/repair	foreign	FALSE
## 2	1	48	5951		2	22 poor	goods/repair	foreign	FALSE
## 3	0	12	2096		2	49 terrible	edu	foreign	FALSE
## 4	0	42	7882		2	45 poor	goods/repair	foreign	FALSE
## 5	1	24	4870		3	53 poor	newcar	foreign	FALSE
## 6	0	36	9055		2	35 poor	edu	foreign	FALSE

```
dim(credit)
```

```
## [1] 1000 9
```

# Logit Demo

```
mylogit <- glm(Default~duration + amount + installment + age
               + factor(history) + factor(purpose) + factor(foreign) + factor(rent),
               data = train, family = "binomial")
summary(mylogit)
```

```
##
## ...
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -3.285e-01  5.597e-01  -0.587  0.557264
## duration      1.625e-02  9.538e-03   1.704  0.088369 .
## amount        1.518e-04  4.325e-05   3.511  0.000447 ***
## installment    3.335e-01  9.216e-02   3.619  0.000296 ***
## age           -1.762e-02  8.851e-03  -1.990  0.046554 *
## factor(history)poor -1.212e+00  3.126e-01  -3.876  0.000106 ***
## factor(history)terrible -1.989e+00  3.552e-01  -5.598  2.17e-08 ***
## factor(purpose)usedcar -1.813e+00  4.067e-01  -4.459  8.23e-06 ***
## factor(purpose)goods/repair -7.163e-01  2.254e-01  -3.177  0.001486 **
## factor(purpose)edu     1.207e-01  3.858e-01   0.313  0.754450
## factor(purpose)biz     -9.862e-01  3.440e-01  -2.867  0.004147 **
## factor(foreign)german -2.057e+00  8.213e-01  -2.505  0.012254 *
## factor(rent)TRUE      7.554e-01  2.355e-01   3.208  0.001337 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## ...
```

# Logit Demo

```
test$phat<- predict(mylogit, test, type="response")
test$Default_hat<-ifelse(test$phat>.5,1,0)
with(test,prop.table(table(Default,Default_hat)))
```

##		Default_hat	
##	Default	0	1
##	0	0.63666667	0.06666667
##	1	0.22666667	0.07000000



# LDA: Demo

$$p(Y = 1) = \frac{\sum_{i=1}^n 1[Y_i = 1]}{N} \quad (29)$$

```
p1<-sum(train$Default)/dim(train)[1]
p1
```

```
## [1] 0.3014286
```

$$\hat{\mu}_k = \frac{1}{n_k} \sum_{i:y_i=k} x_i \quad (30)$$

```
mu1<-mean(train$duration[train$Default==1])
mu1
```

```
## [1] 24.78673
```

```
mu0<-mean(train$duration[train$Default==0])
mu0
```

```
## [1] 19.79346
```

# LDA: Demo

$$\hat{\sigma}^2 = \frac{1}{N - K} \sum_{k=1}^K \sum_{i: y_i = k} (x_i - \hat{\mu}_k)^2 \quad (31)$$

```
g1<-sum((train$duration[train$Default==1]-mu1)^2)
g0<-sum((train$duration[train$Default==0]-mu0)^2)

sigma<-sqrt((g1+g0)/(dim(train)[1]-2))
```

$$\hat{f}_k \sim N(\hat{\mu}_k, \hat{\sigma}) \quad (32)$$

```
f1<-dnorm(test$duration,mean=mu1,sd=sigma)
f0<-dnorm(test$duration,mean=mu0,sd=sigma)
```

# LDA: Demo

```
library("MASS")      # LDA
lda_simple <- lda(Default~duration, data = train)
lda_simple_pred<-predict(lda_simple,test)
names(lda_simple_pred)
```

```
## [1] "class"      "posterior" "x"
```

```
posteriors<-data.frame(lda_simple_pred$posterior)
posteriors$hand<-f1*p1/(f1*p1+f0*(1-p1))
head(posteriors)
```


```
##           X0           X1          hand
## 1  0.8013656 0.1986344 0.1986344
## 3  0.7668614 0.2331386 0.2331386
## 14 0.6861792 0.3138208 0.3138208
## 16 0.6861792 0.3138208 0.3138208
## 28 0.7668614 0.2331386 0.2331386
## 33 0.7283950 0.2716050 0.2716050
```

# Roc Demo

```
library("ROCR") #Roc


mylogit <- glm(Default~duration + amount + installment + age
               + factor(history) + factor(purpose) + factor(foreign) + factor(rent),
               data = train, family = "binomial")

test$phat<- predict(mylogit, test, type="response")
pred <- prediction(test$phat, test$Default)
roc_ROCR <- performance(pred,"tpr","fpr")
plot(roc_ROCR, main = "ROC curve", colorize = T)
abline(a = 0, b = 1)
```





# Roc Demo



```
auc_ROCR <- performance(pred, measure = "auc")  
auc_ROCR@y.values[[1]]
```

```
## [1] 0.714415
```

# Roc Demo


```
mylda <- lda(Default~duration + amount + installment + age , data = train)
mylda
```

```
## Call:
## lda(Default ~ duration + amount + installment + age, data = train)
##
## Prior probabilities of groups:
##      0      1
## 0.6985714 0.3014286
##
## Group means:
##   duration  amount installment    age
## 0 19.79346 3062.888   2.885481 36.40900
## 1 24.78673 4057.791   3.109005 33.85782
##
## Coefficients of linear discriminants:
##               LD1
## duration    0.0296041361
## amount      0.0002055164
## installment 0.4821242957
## age        -0.0386710882
```

# Roc Demo


```
phat_mylda <- predict(mylda, test, type="response")
pred_mylda <- prediction(phat_mylda$posterior[,2], test$Default)

roc_mylda <- performance(pred_mylda, "tpr", "fpr")
plot(roc_mylda, main = "ROC curve", colorize = T)
abline(a = 0, b = 1)
```



# Roc Demo

```
plot(roc_ROCR, main = "ROC curve", colorize = FALSE, col="red")  
plot(roc_mylda, add=TRUE, colorize = FALSE, col="blue")  
abline(a = 0, b = 1)
```



# Roc Demo

## ► Area under the curve (AUC)

```
auc_ROCR <- performance(pred, measure = "auc")  
auc_ROCR_lda_simple <- performance(pred_mylda, measure = "auc")  
auc_ROCR@y.values[[1]]
```

```
## [1] 0.714415
```

```
auc_ROCR_lda_simple@y.values[[1]]
```

```
## [1] 0.6291602
```