

Lecture 18: Classification

Big Data and Machine Learning for Applied Economics
Econ 4676

Ignacio Sarmiento-Barbieri

Universidad de los Andes

October 19, 2021

Agenda

- 1 Recap: Regularization and Lasso for Causality
- 2 Classification
 - K-Nearest Neighbors
 - Logit
 - Linear Discriminant Analysis
- 3 Misclassification Rates
 - ROC curve
- 4 Review & Next Steps
- 5 Further Readings
- 6 Demos in R
 - KNN
 - Logit
 - LDA
 - ROC

Elastic Net

► Naive Elastic Net

$$\min_{\beta} NEL(\beta) = \sum_{i=1}^n (y_i - x'_i \beta)^2 + \lambda_1 \sum_{s=2}^p |\beta_s| + \lambda_2 \sum_{s=2}^p \beta_s^2 \quad (1)$$

- Elastic Net: rescaled version. Double Shrinkage introduces “too” much bias, *final* version “corrects” for this

$$\hat{\beta}_{EN} = \frac{1}{\sqrt{1 + \lambda_2}} \hat{\beta}_{naive EN} \quad (2)$$

Lasso for Causality

Inference with Selection among Many Controls

$$y_i = \alpha D_i + X_i' \theta_y + r_{yi} + \zeta_i \quad (3)$$

- ▶ We apply variable selection methods to each of the two reduced form equations and then use all of the selected controls in estimation of α .
- ▶ We select
 - 1 A set of variables that are useful for predicting y_i , say X_{yi} , and
 - 2 A set of variables that are useful for predicting D_i , say X_{di} .
- ▶ We then estimate α by ordinary least squares regression of y_i on d_i and the union of the variables selected for predicting y_i and D_i , contained in X_{yi} and X_{di} .
- ▶ We thus make sure we use variables that are important for either of the two predictive relationships to guard against OVB

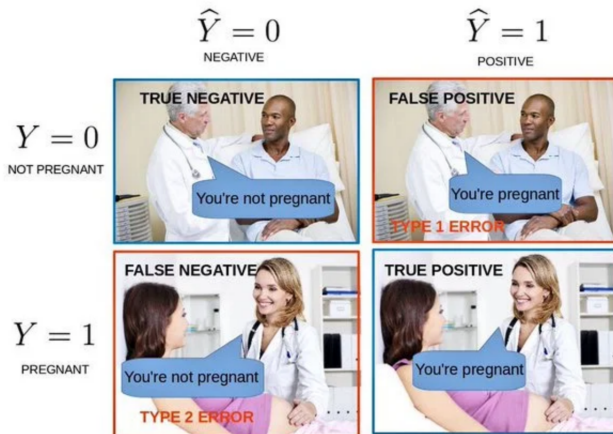
Classification

Classification: Motivation

- ▶ Admit a student to *PEG* based on their grades and LoR
- ▶ Give a credit, based on credit history, demographics?
- ▶ Classifying emails: spam, personal, social based on email contents
- ▶ Aim is to classify y based on X 's
- ▶ y can be
 - ▶ qualitative (e.g., spam, personal, social)
 - ▶ Not necessarily ordered
 - ▶ Not necessarily two categories, but will start with the binary case

Motivation

- ▶ Two states of nature $y \rightarrow n \in \{0,1\}$
- ▶ Two actions $(\hat{y}) \rightarrow a \in \{0,1\}$



Source: <https://dzone.com/articles/understanding-the-confusion-matrix>

Probability, Cost, and Classification

- ▶ Two states of nature $y \rightarrow n \in \{0, 1\}$
- ▶ Two actions $(\hat{y}) \rightarrow a \in \{0, 1\}$
- ▶ Probabilities
 - ▶ $p = Pr(y = 1|X)$
 - ▶ $1 - p = Pr(y = 0|X)$
- ▶ Loss: $L(a,)$, penalizes being in bin (a, n)
- ▶ Risk: expected loss of taking action a

Probability, Cost, and Classification

- Risk: expected loss of taking action a

$$E[L(a, n)] = \sum_n p_n L(a, n) \quad (4)$$

$$R(a) = (1 - p)L(a, 0) + pL(a, 1)$$

- The objective is the same as before: minimize the risk
- We have to define $L(a, n)$

Probability, Cost, and Classification

- ▶ Which action do we choose?

Probability, Cost, and Classification

- ▶ Which action do we choose?
- ▶ We can compare the risk of each action
- ▶ We are going to choose to take action 1 when the risk is lower:

$$\begin{aligned} R(1) &< R(0) \\ 1 - p &< p \\ p &> \frac{1}{2} \end{aligned} \tag{5}$$

- ▶ This is known as the Bayes Classifier, choose the estate that minimizes the risk


Probability, Cost, and Classification

- ▶ Under a 0-1 penalty the problem boils down to finding $p = Pr(y = 1|X)$
- ▶ We then predict 1 if $p > 0.5$ and 0 otherwise (Bayes classifier)
- ▶ We can think 3 ways of finding this probability in binary cases
 - ▶ K-Nearest Neighbors
 - ▶ Logistic
 - ▶ LDA
- ▶ Why not $p = X\beta$?

K-Nearest Neighbors

K-Nearest Neighbors

- K nearest neighbor (K-NN) algorithm predicts class \hat{y} for x by asking *What is the most common class for observations around x ?*



Source: Taddy (2019)

Lecture 18

K-Nearest Neighbors

- ▶ K nearest neighbor (K-NN) algorithm predicts class \hat{y} for x by asking *What is the most common class for observations around x ?*
- ▶ Algorithm: given an input vector x_f where you would like to predict the class label
 - ▶ Find the K nearest neighbors in the dataset of labeled observations, $x_i, y_{i=1}^n$, the most common distance is the Euclidean distance (units):

$$d(x_i, x_f) = \sqrt{\sum_{j=1}^p (x_{ij} - x_{fj})^2} \quad (6)$$

- ▶ This yields a set of the K nearest observations with labels:

$$[x_{i1}, y_{i1}], \dots, [x_{iK}, y_{iK}] \quad (7)$$

- ▶ The predicted class of x_f is the most common class in this set

$$\hat{y}_f = \text{mode}\{y_{i1}, \dots, y_{iK}\} \quad (8)$$

K-Nearest Neighbors


- ▶ There are some major problems with practical applications
 - ▶ Knn predictions are unstable as a function of K

$$K = 1 \implies \hat{p}(\text{white}) = 0$$

$$K = 2 \implies \hat{p}(\text{white}) = 1/2$$

$$K = 3 \implies \hat{p}(\text{white}) = 2/3$$

$$K = 4 \implies \hat{p}(\text{white}) = 1/2$$



Source: Taddy (2019)

K-Nearest Neighbors

- ▶ In this case
 - ▶ 1-Knn manages 70% accuracy
 - ▶ 5-Knn manages 60% accuracy
 - ▶ H.W. try for different seed, (Taddy's is 80% and 70%)
- ▶ There are some major problems with practical implications
 - ▶ Knn predictions are unstable as a function of K
 - ▶ This instability of prediction makes it hard to choose the optimal K and cross validation doesn't work well for KNN
 - ▶ Since prediction for each new x requires a computationally intensive counting, KNN is too expensive to be useful in most big data settings.
 - ▶ KNN is a good idea, but too crude to be useful in practice

Logit

Logit


We have a conditional probability

$$\Pr(y = 1|X) = f(X'\beta) \quad (9)$$

Logistic regression uses a *logit* (sigmoid, softmax) link function

$$\log \left(\frac{p(y = 1|X)}{1 - p(y = 1|X)} \right) = \beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k \quad (10)$$

Logit



Source: Taddy (2019)

Logit

We have a conditional probability

$$Pr(y = 1|X) = f(X'\beta) \quad (11)$$

Can recover predictions:

$$p(y = 1|X) = \frac{e^{X'\beta}}{1 + e^{X'\beta}} = \frac{\exp(\beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k)}{1 + \exp(\beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k)} \quad (12)$$

Linear Discriminant Analysis

Linear Discriminant Analysis

Reverend Bayes to the rescue: Bayes Theorem

$$p(y = 1|X) = \frac{f(X|y = 1)p(y = 1)}{m(X)} \quad (13)$$

with $m(X)$ is the marginal distribution of X , i.e.

$$m(X) = \int f(X|y = 1)p(y = 1)dy \quad (14)$$

Recall that there are two states of nature $y \rightarrow i \in \{0, 1\}$

$$\begin{aligned} m(X) &= f(X|y = 1)p(y = 1) + f(X|y = 0)p(y = 0) \\ &= f(X|y = 1)p(y = 1) + f(X|y = 0)(1 - p(y = 1)) \end{aligned} \quad (15)$$

Linear Discriminant Analysis

- ▶ This is basically an empirical Bayes approach
- ▶ We need to estimate $f(X|y = 1), f(X|y = 0)$ and $p(y = 1)$
 - ▶ Let's start by estimating $p(y = 1)$. We've done this before

$$p(y = 1) = \frac{\sum_{i=1}^n 1[y_i = 1]}{N} \quad (16)$$

- ▶ Next $f(X|y = j)$ with $j = 0, 1$.
 - ▶ if we assume one predictor and $X|y \sim N(\mu_j, \sigma_j)$
 - ▶ the problem boils down to estimating μ_j, σ_j
 - ▶ LDA makes it simpler, assumes $\sigma_j = \sigma \forall j$
 - ▶ then partition the sample in two $y = 0$ and $y = 1$, estimate the moments and get $\hat{f}(X|y = j)$
- ▶ Plug everything into the Bayes Rule and you're done

Linear Discriminant Analysis

Extensions

- ▶ If we have k predictors?
- ▶ then $X|y \sim NM(\mu, \Sigma)$

$$f(X|y = j) = \frac{1}{(2\pi)^{k/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_j)' \Sigma_j (x - \mu_j)\right) \quad (17)$$

- ▶ μ_j is the vector of the sample means in each partition $j = 0, 1$
- ▶ Σ_j is the matrix of variance and covariances of each partition $j = 0, 1$
- ▶ Can we lift normality?

Linear Discriminant Analysis

- ▶ Why is it call linear?
- ▶ Note

$$p > \frac{1}{2} \iff \ln\left(\frac{p}{(1-p)}\right) \quad (18)$$

- ▶ Logit with one predictor

$$\beta_1 + \beta_2 X \quad (19)$$

- ▶ Classification: in the probability of space
- ▶ Discrimination: in the space of X
- ▶ $\beta_1 + \beta_2 X$ is the discrimination function for logit (it is lineal)

Linear Discriminant Analysis

- ▶ LDA?
- ▶ One predictor with $\sigma_0 = \sigma_1$ (equal variance)

$$p(y = 1|X) = \frac{f(X|y = 1)p(y = 1)}{f(X|y = 1)p(y = 1) + f(X|y = 0)(1 - p(y = 1))} \quad (20)$$

- ▶ Then under the equal variance assumption

$$\frac{p(y = 1|X)}{1 - p(y = 1|X)} = \frac{f(X|y = 1)p(y = 1)}{f(X|y = 0)(1 - p(y = 1))} \quad (21)$$

$$= \frac{p(y = 1)\exp(-(x - \mu_1)^2)}{(1 - p(y = 1))\exp(-(x - \mu_0)^2)} \quad (22)$$

Linear Discriminant Analysis

► Taking logs

$$\log \left(\frac{p(y = 1|X)}{1 - p(y = 1|X)} \right) = \log \left(\frac{p(y = 1)}{(1 - p(y = 1))} + (x - \mu_1)^2 - (x - \mu_0)^2 \right) \quad (23)$$

$$= \log \left(\frac{p(y = 1)}{(1 - p(y = 1))} + \mu_1^2 - \mu_0^2 - 2(\mu_1 - \mu_0)x \right) \quad (24)$$


$$= \gamma_1 + \gamma_2 X \quad (25)$$

- under the assumption of equal variance the discrimination function is lineal
- Note: logit estimates γ_1 and γ_2

Misclassification Rates


Misclassification Rates

- Predicted probabilities from Logit model



Misclassification Rates

- ▶ A classification rule, or cutoff, is the probability p at which you predict
 - ▶ $\hat{y}_i = 0$ if $p_i < p$
 - ▶ $\hat{y}_i = 1$ if $p_i \geq p$
- ▶ Measures of performance
 - ▶ *1-Specificity*: False Positive Rate, Type I error
 - ▶ *Sensitivity*: True Positive Rate, power, (1-Type II error)




ROC

- ▶ ROC curve: Receiver operating characteristic curve
- ▶ ROC curve illustrates the trade-off of the classification rule
- ▶ Gives us the ability
 - ▶ Measure the predictive capacity of our model
 - ▶ Compare between models
- ▶ Some definitions
 - ▶ $P = \sum y_i$ positives
 - ▶ $N = \sum (1 - y_i)$ negatives
 - ▶ $T = P + N$ all observations
 - ▶ True Positives: $TP = \sum \hat{y}_i y_i$, True Positive Rate = $\frac{TP}{P}$
 - ▶ False Positives: $FP = \sum \hat{y}_i (1 - y_i)$, False Positive Rate = $\frac{FP}{N}$

ROC

- ▶ Binary Classifier: $\hat{y}_i = 1[p_i > c]$, $c \in [0, 1]$
- ▶ Bayes fixes $c = 0.5$
- ▶ Ideally $TPR = 1$ and $FPR = 0$
- ▶ ROC curve give us the locus of all possible TPR and FPR for all possible $c \in [0, 1]$



► ROC Properties

► Has positive slope

► In $(0,0)$, $c = 1$. When $c \downarrow$, $TP \uparrow$ and $FP \uparrow$. Then

$$TPR = \sum \frac{\hat{y}_i y_i}{P} \quad FPR = \sum \frac{\hat{y}_i (1 - y_i)}{T - P} \quad (26)$$

► Is easy to show

$$TPR = \frac{\sum \hat{y}_i}{P} - \frac{T - P}{P} FPR \quad (27)$$

► ROC is the locus of all possible TPR and FPR for all possible $c \in [0, 1]$

$$TPR = \frac{\sum \hat{y}_i(c)}{P} - \frac{T - P}{P} FPR(c) \quad (28)$$

ROC: Summary

- ▶ Ideal ROC curve
- ▶ AUC: area under the curve, is like an R^2
- ▶ Help us compare between classifiers
- ▶ Dominated classifiers?
- ▶ Which c ? Choose a max FPR

Review & Next Steps

- ▶ Review Classification:
 - ▶ KNN
 - ▶ Intuitive
 - ▶ Not very useful in practice, curse of dimensionality
 - ▶ Logit
 - ▶ Linear Discriminant Analysis
 - ▶ Misclassification Rates: ROC curve
 - ▶ QDA?
 - ▶ Multiple Classes?
- ▶ Next class: Problem Sets, Text Data!
- ▶ Questions? Questions about software?

Further Readings

- ▶ Friedman, J., Hastie, T., & Tibshirani, R. (2001). The elements of statistical learning (Vol. 1, No. 10). New York: Springer series in statistics.
- ▶ James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). An introduction to statistical learning (Vol. 112, p. 18). New York: springer.
- ▶ Kuhn, M. (2012). The caret package. R Foundation for Statistical Computing, Vienna, Austria. <https://topepo.github.io/caret/index.html>
- ▶ Taddy, M. (2019). Business data science: Combining machine learning and economics to optimize, automate, and accelerate business decisions. McGraw Hill Professional.
- ▶ Zou, H. y Hastie, T., 2005, Regularization and variable selection via the elastic net, Journal of the Royal Statistical Society, 67, 2, 301-320.


K-Nearest Neighbors

```
#Load the required packages
library("class") #for KNN
library("MASS") # a library of example datasets
#Read the data
data(fgl) ## loads the data into R; see help(fgl)
str(fgl)
```

```
## 'data.frame':    214 obs. of  10 variables:
## $ RI : num  3.01 -0.39 -1.82 -0.34 -0.58 ...
## $ Na : num  13.6 13.9 13.5 13.2 13.3 ...
## $ Mg : num  4.49 3.6 3.55 3.69 3.62 3.61 3.6 3.61 3.58 3.6 ...
## $ Al : num  1.1 1.36 1.54 1.29 1.24 1.62 1.14 1.05 1.37 1.36 ...
## $ Si : num  71.8 72.7 73 72.6 73.1 ...
## $ K : num  0.06 0.48 0.39 0.57 0.55 0.64 0.58 0.57 0.56 0.57 ...
## $ Ca : num  8.75 7.83 7.78 8.22 8.07 8.07 8.17 8.24 8.3 8.4 ...
## $ Ba : num  0 0 0 0 0 0 0 0 0 0 ...
## $ Fe : num  0 0 0 0 0 0.26 0 0 0 0.11 ...
## $ type: Factor w/ 6 levels "WinF","WinNF",...: 1 1 1 1 1 1 1 1 1 1 ...
```

Refractive index and chemical composition for six possible glass types: float glass window (WinF), nonfloat glass window (WinNF), vehicle window (Veh), container (Con), tableware (Tabl), vehicle headlamp (Head)

K-Nearest Neighbors



K-Nearest Neighbors

- ▶ Units matter
 - ▶ Since distance is measured on the raw x values, units matter.
 - ▶ As we did for regularization, we will standardized observations.
 - ▶ R `scale` function does this, i.e., convert columns to mean-zero sd-one

```
x <- scale(fgl[,1:9]) # column 10 is the class label  
apply(x,2,sd) # see ?apply
```

```
## RI Na Mg Al Si K Ca Ba Fe  
## 1 1 1 1 1 1 1 1 1
```


K-Nearest Neighbors

► Before running Knn

- Make sure you have numeric matrices of training data x values, with labels y
- Also need to provide new *test* values where you would like to predict
- Note that there's no model to fit, Knn, just counts neighbors for each observation in test

```
set.seed(1010101)
test <- sample(1:214,10)
nearest1 <- knn(train=x[-test,], test=x[test,], cl=fgl$type[-test], k=1)
nearest5 <- knn(train=x[-test,], test=x[test,], cl=fgl$type[-test], k=5)
data.frame(fgl$type[test],nearest1,nearest5)
```

```
##      fgl.type.test. nearest1 nearest5
## 1           WinF         WinF   WinNF
## 2            Head         Head    Head
## 3          WinNF      WinNF   WinNF
## 4           WinF         WinF    WinF
## 5          WinNF      WinNF   WinNF
## 6          WinNF      WinNF   WinNF
## 7            Head          Con     Con
## 8            Head      WinNF   WinNF
## 9          WinNF      WinNF   WinNF
## 10         WinNF         WinF    WinF
```

Logit Demo

```
set.seed(101010) #sets a seed
credit<-readRDS("credit_class.rds")
#70% train
indic<-sample(1:nrow(credit),floor(.7*nrow(credit)))
#Partition the sample
train<-credit[indic,]
test<-credit[-indic,]
head(credit)
```

```
##   Default duration amount installment age  history      purpose foreign  rent
## 1      0         6    1169           4  67 terrible goods/repair foreign FALSE
## 2      1        48    5951           2  22      poor goods/repair foreign FALSE
## 3      0        12    2096           2  49 terrible          edu foreign FALSE
## 4      0        42    7882           2  45      poor goods/repair foreign FALSE
## 5      1        24    4870           3  53      poor    newcar foreign FALSE
## 6      0        36    9055           2  35      poor          edu foreign FALSE
```

```
dim(credit)
```

```
## [1] 1000    9
```

Logit Demo

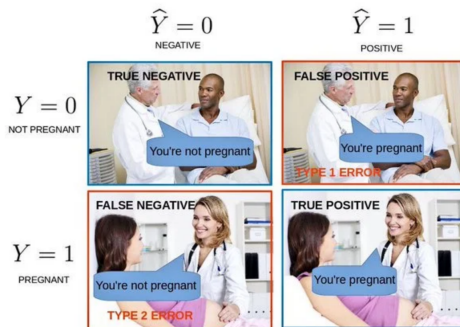
```
mylogit <- glm(Default~duration + amount + installment + age
               + factor(history) + factor(purpose) + factor(foreign) + factor(rent),
               data = train, family = "binomial")
summary(mylogit)
```

```
##
## ...
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.285e-01  5.597e-01  -0.587  0.557264
## duration        1.625e-02  9.538e-03   1.704  0.088369 .
## amount          1.518e-04  4.325e-05   3.511  0.000447 ***
## installment     3.335e-01  9.216e-02   3.619  0.000296 ***
## age            -1.762e-02  8.851e-03  -1.990  0.046554 *
## factor(history)poor -1.212e+00  3.126e-01  -3.876  0.000106 ***
## factor(history)terrible -1.989e+00  3.552e-01  -5.598  2.17e-08 ***
## factor(purpose)usedcar -1.813e+00  4.067e-01  -4.459  8.23e-06 ***
## factor(purpose)goods/repair -7.163e-01  2.254e-01  -3.177  0.001486 **
## factor(purpose)edu      1.207e-01  3.858e-01   0.313  0.754450
## factor(purpose)biz      -9.862e-01  3.440e-01  -2.867  0.004147 **
## factor(foreign)german -2.057e+00  8.213e-01  -2.505  0.012254 *
## factor(rent)TRUE       7.554e-01  2.355e-01   3.208  0.001337 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## ...
```

Logit Demo

```
test$phat<- predict(mylogit, test, type="response")
test$Default_hat<-ifelse(test$phat>.5,1,0)
with(test,prop.table(table(Default,Default_hat)))
```

```
##           Default_hat
## Default           0           1
##      0 0.63666667 0.06666667
##      1 0.22666667 0.07000000
```



LDA: Demo

$$p(y = 1) = \frac{\sum_{i=1}^n 1[y_i = 1]}{N} \quad (29)$$

```
p1<-sum(train$Default)/dim(train)[1]
p1
```

```
## [1] 0.3014286
```

$$\hat{\mu}_k = \frac{1}{n_k} \sum_{i:y_i=k} x_i \quad (30)$$

```
mu1<-mean(train$duration[train$Default==1])
mu1
```

```
## [1] 24.78673
```

```
mu0<-mean(train$duration[train$Default==0])
mu0
```

```
## [1] 19.79346
```

LDA: Demo

$$\hat{\sigma}^2 = \frac{1}{N - K} \sum_{k=1}^K \sum_{i: y_i = k} (x_i - \hat{\mu}_k)^2 \quad (31)$$

```
g1<-sum((train$duration[train$Default==1]-mu1)^2)
g0<-sum((train$duration[train$Default==0]-mu0)^2)

sigma<-sqrt((g1+g0)/(dim(train)[1]-2))
```

$$\hat{f}_k \sim N(\hat{\mu}_k, \hat{\sigma}) \quad (32)$$

```
f1<-dnorm(test$duration,mean=mu1,sd=sigma)
f0<-dnorm(test$duration,mean=mu0,sd=sigma)
```

LDA: Demo

```
library("MASS")      # LDA
lda_simple <- lda(Default~duration, data = train)
lda_simple_pred<-predict(lda_simple,test)
names(lda_simple_pred)
```

```
## [1] "class"      "posterior" "x"
```

```
posteriors<-data.frame(lda_simple_pred$posterior)
posteriors$hand<-f1*p1/(f1*p1+f0*(1-p1))
head(posteriors)
```


```
##           X0           X1           hand
## 1  0.8013656 0.1986344 0.1986344
## 3  0.7668614 0.2331386 0.2331386
## 14 0.6861792 0.3138208 0.3138208
## 16 0.6861792 0.3138208 0.3138208
## 28 0.7668614 0.2331386 0.2331386
## 33 0.7283950 0.2716050 0.2716050
```

Roc Demo


```
library("ROCR") #Roc

mylogit <- glm(Default~duration + amount + installment + age
               + factor(history) + factor(purpose) + factor(foreign) + factor(rent),
               data = train, family = "binomial")

test$phat<- predict(mylogit, test, type="response")
pred <- prediction(test$phat, test$Default)
roc_ROCR <- performance(pred,"tpr","fpr")
plot(roc_ROCR, main = "ROC curve", colorize = T)
abline(a = 0, b = 1)
```



Roc Demo



```
auc_ROCR <- performance(pred, measure = "auc")  
auc_ROCR@y.values[[1]]
```

```
## [1] 0.714415
```

Roc Demo


```
mylda <- lda(Default~duration + amount + installment + age , data = train)
mylda
```

```
## Call:
## lda(Default ~ duration + amount + installment + age, data = train)
##
## Prior probabilities of groups:
##      0      1
## 0.6985714 0.3014286
##
## Group means:
##      duration  amount installment      age
## 0 19.79346 3062.888    2.885481 36.40900
## 1 24.78673 4057.791    3.109005 33.85782
##
## Coefficients of linear discriminants:
##                LD1
## duration      0.0296041361
## amount        0.0002055164
## installment   0.4821242957
## age          -0.0386710882
```

Roc Demo


```
phat_mylda <- predict(mylda, test, type="response")
pred_mylda <- prediction(phat_mylda$posterior[,2], test$Default)

roc_mylda <- performance(pred_mylda, "tpr", "fpr")
plot(roc_mylda, main = "ROC curve", colorize = T)
abline(a = 0, b = 1)
```



Roc Demo

```
plot(roc_ROCR, main = "ROC curve", colorize = FALSE, col="red")  
plot(roc_mylda, add=TRUE, colorize = FALSE, col="blue")  
abline(a = 0, b = 1)
```



Roc Demo

► Area under the curve (AUC)

```
auc_ROCR <- performance(pred, measure = "auc")  
auc_ROCR_lda_simple <- performance(pred_mylda, measure = "auc")  
auc_ROCR@y.values[[1]]
```

```
## [1] 0.714415
```

```
auc_ROCR_lda_simple@y.values[[1]]
```

```
## [1] 0.6291602
```